

Part 3

Memory Management

3.1 No Memory Abstraction

3.2 Memory Abstraction

3.3 Virtual memory

3.4 Page Replacement Algorithms

3.5 Design Issues For Paging Systems

3.6 Implementation issues

3.7 Segmentation

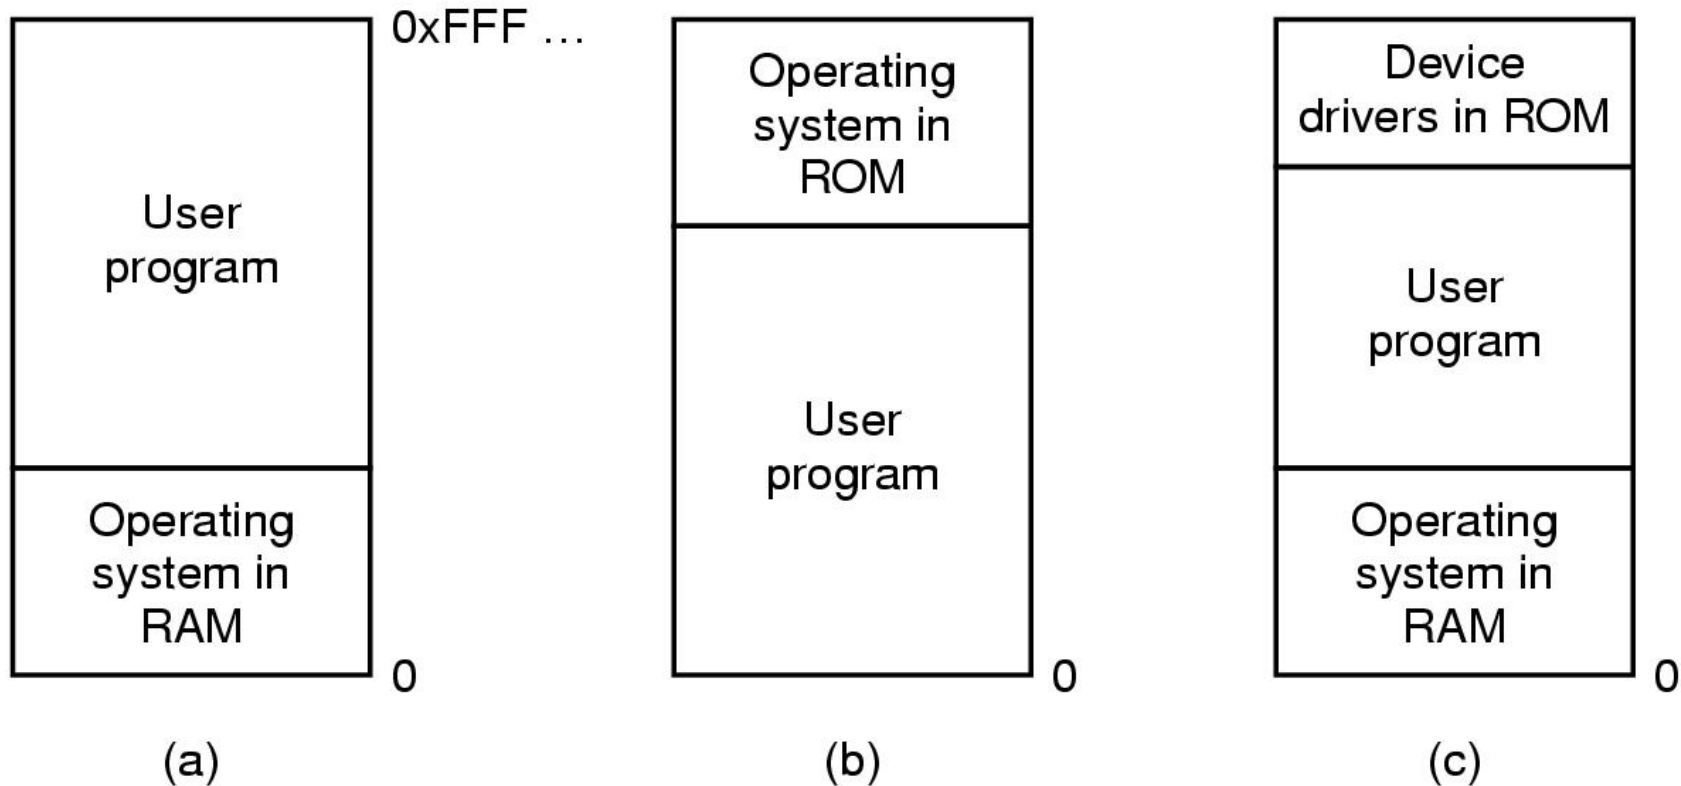
Memory Management

- Ideally programmers want memory that is
 - large
 - fast
 - non volatile
- Memory hierarchy
 - small amount of fast, expensive memory – cache
 - some medium-speed, medium price main memory
 - gigabytes of slow, cheap disk storage
- Memory manager handles the memory hierarchy

3.1 Memory Abstraction

No Memory Abstraction

Monoprogramming without Swapping or Paging



Three simple ways of organizing memory

- an operating system with one user process

No Memory Abstraction

Multiple Programs Without Memory Abstraction

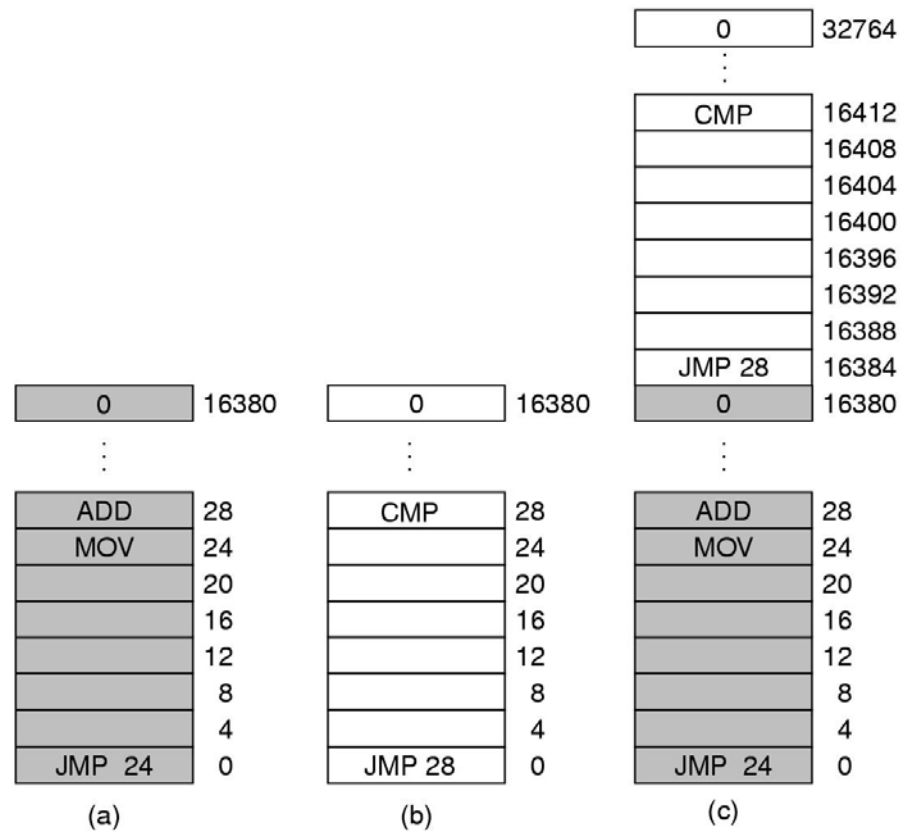
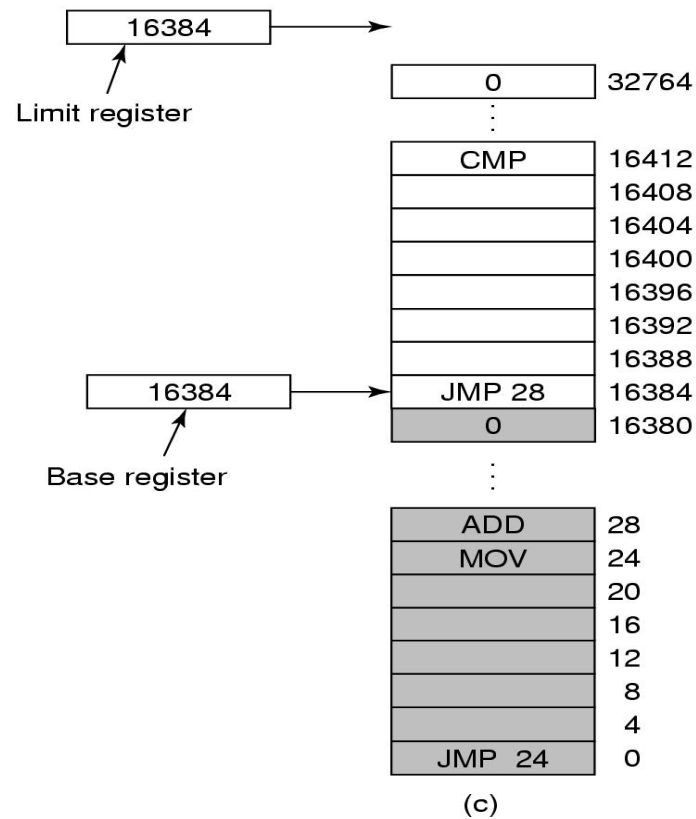


Illustration of the relocation problem.

3.2 Memory Abstraction

Memory Abstraction


Base and Limit Registers



Base and limit registers can be used to give each process a separate address space.

Memory Abstraction

Dynamic relocation using a relocation register



Memory Abstraction

Relocation and Protection

- Cannot be sure where program will be loaded in memory
 - address locations of variables, code routines cannot be absolute
 - must keep a program out of other processes' partitions
- Use base and limit values
 - address locations added to base value to map to physical addr
 - address locations larger than limit value is an error

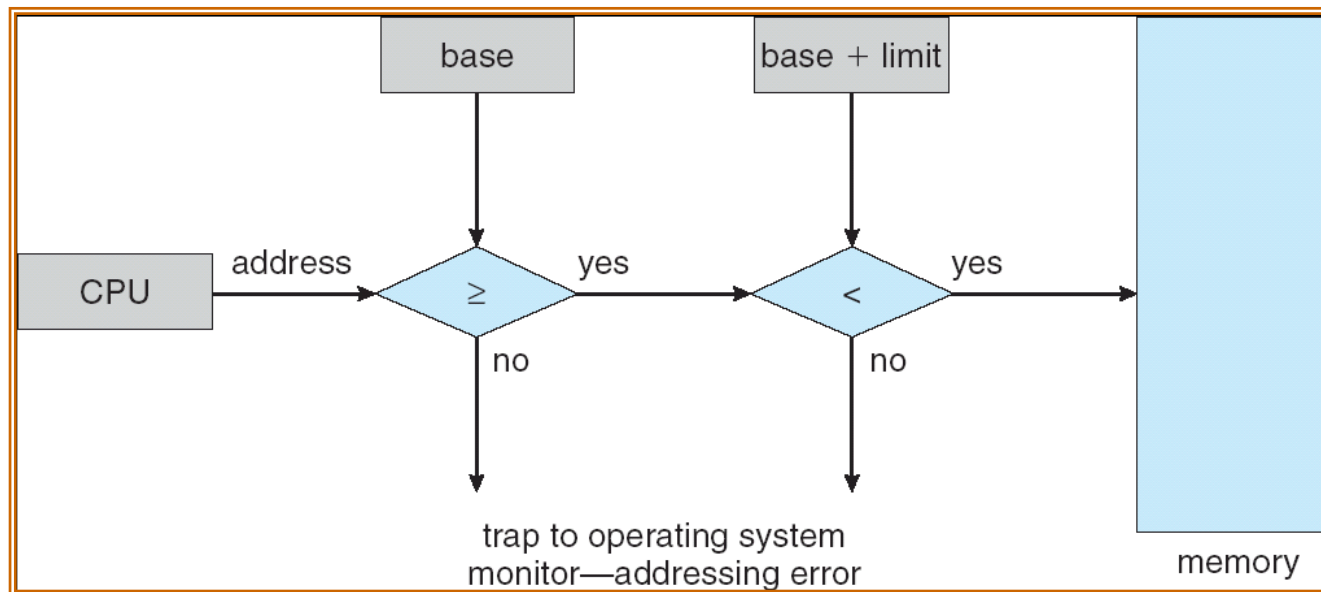
Memory Abstraction

Relocation and Protection

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*

Memory Abstraction


HW address protection with base and limit registers



Memory Abstraction


Swapping (1)

Schematic View of Swapping



Memory Abstraction

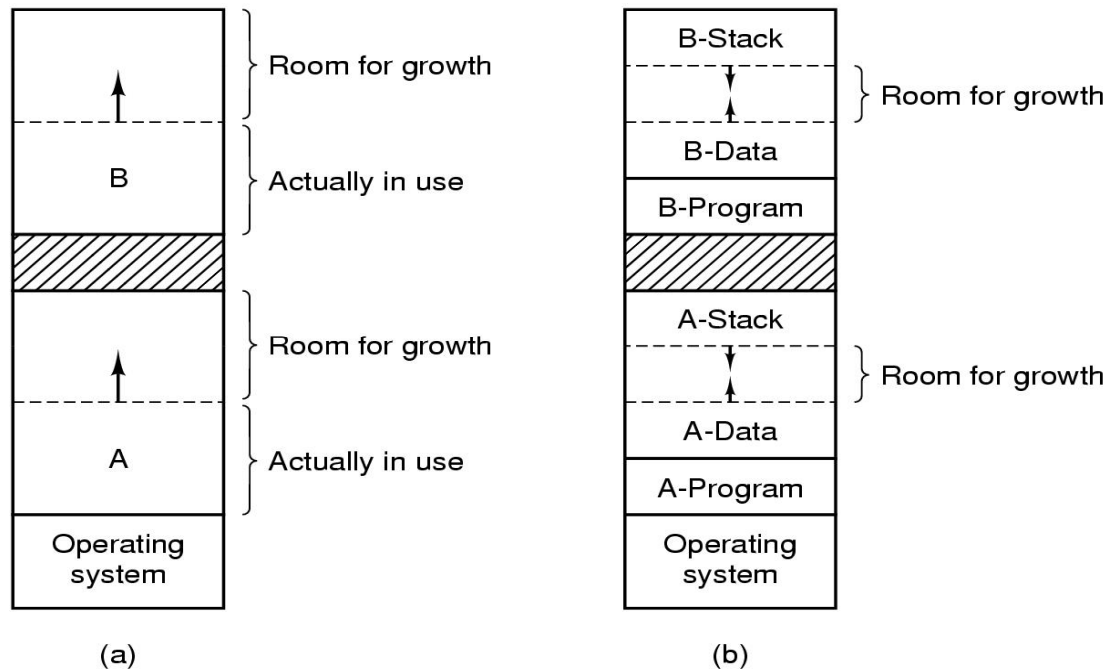
Swapping (2)



- Memory allocation changes as
 - processes come into memory
 - leave memory
- Shaded regions are unused memory
- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used

Memory Abstraction

Swapping (3)



- (a) Allocating space for growing data segment
- (b) Allocating space for growing stack & data segment

Memory Abstraction

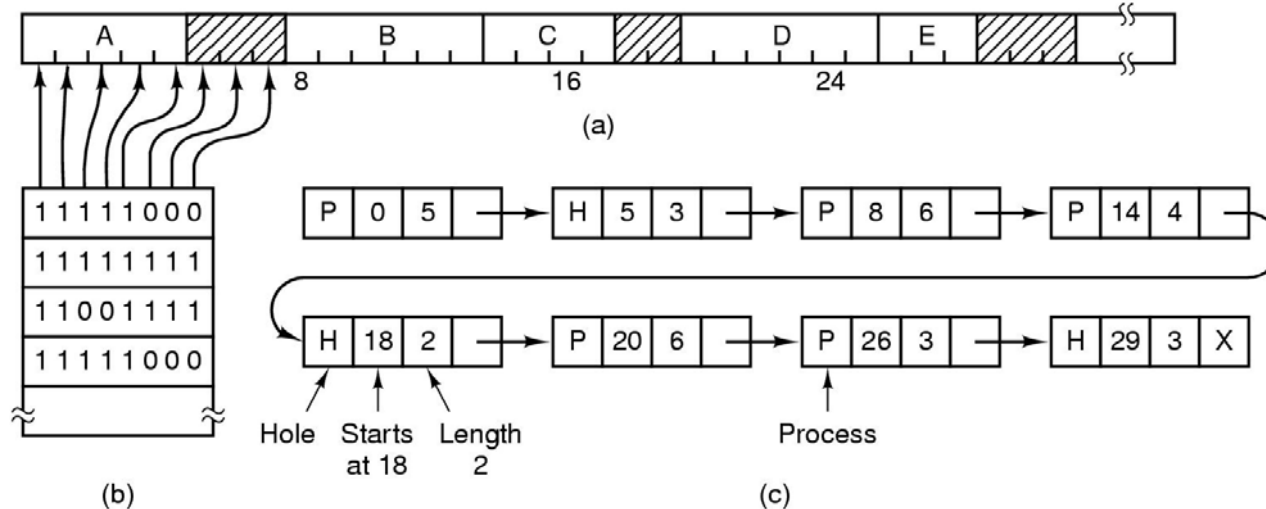
Swapping (4)

Multiple-partition allocation

- Multiple-partition allocation
 - Hole – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)
 - There are two ways to keep track of memory usages
 - Memory Management with Bit Maps
 - Memory Management with Linked Lists

Memory Abstraction Swapping (4)

Multiple-partition allocation




Memory Management with Bit Maps

- (a) Part of memory with 5 processes, 3 holes
 - tick marks show allocation units
 - shaded regions are free
- (b) Corresponding bit map
- (c) Same information as a list

Memory Abstraction

Swapping (5)

Multiple-partition allocation



Four neighbor combinations
for the terminating process, X.

Memory Abstraction

Swapping (6)

Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes


- **First-fit:** Allocate the *first* hole that is big enough
- **Next fit:** Start searching the list from the place where it left off last time
- **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit:** Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization

3.3 Virtual Memory


Virtual Memory

- **Virtual memory** – separation of user logical memory from physical memory.
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
- **Virtual memory can be implemented via:**
 - Demand **paging**
 - Demand **segmentation**

Virtual Memory Paging



Virtual Memory Paging



The position and function of the MMU

Virtual Memory

Paging

- Virtual address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
- Divide physical memory into fixed-sized blocks called **Page frames** (size is power of 2, between 512 bytes and 8,192 bytes)
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size n pages, need to find n free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Internal fragmentation

Virtual Memory

Address Translation Scheme

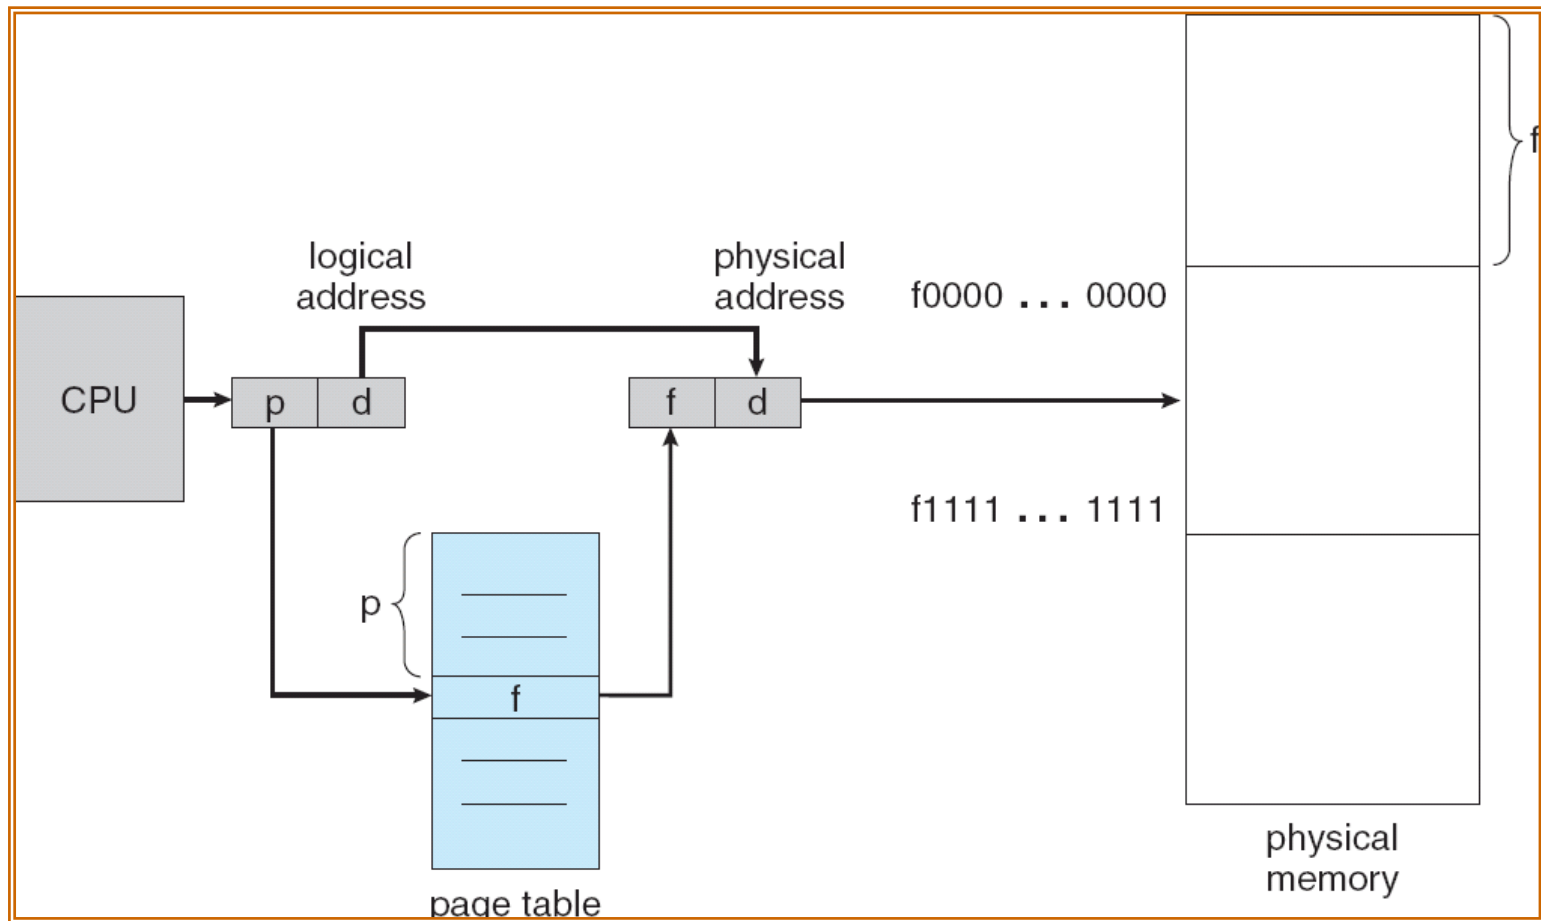
- Address generated by CPU is divided into:
 - **Page number (p)** – used as an index into a *page table* which contains base address of each page in physical memory

page number	page offset
p	d
$m - n$	n

- **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit

Virtual Memory


Paging Hardware



Virtual Memory


Paging: Example

The relation between virtual addresses and physical memory addresses given by page table



Virtual Memory


Page Tables: Example



Internal operation of MMU with 16 4 KB pages

Virtual Memory

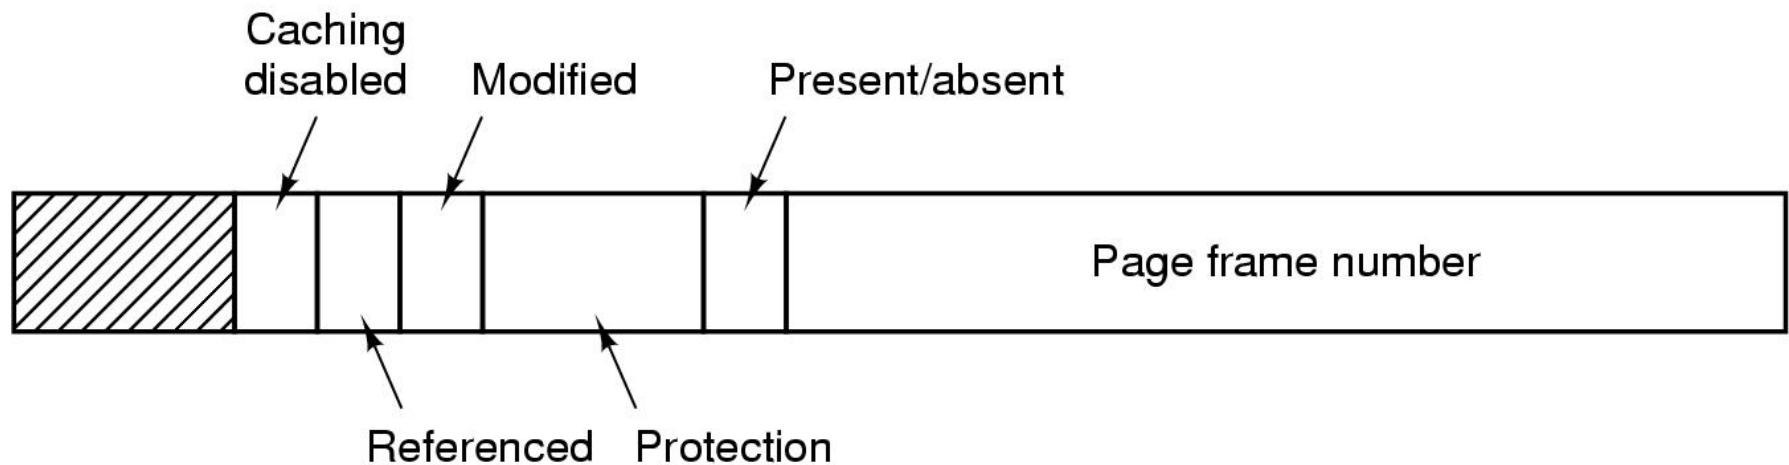
Two-level page tables



- 32 bit address with 2 page table fields
- Two-level page tables

Virtual Memory

Typical page table entry



Typical page table entry

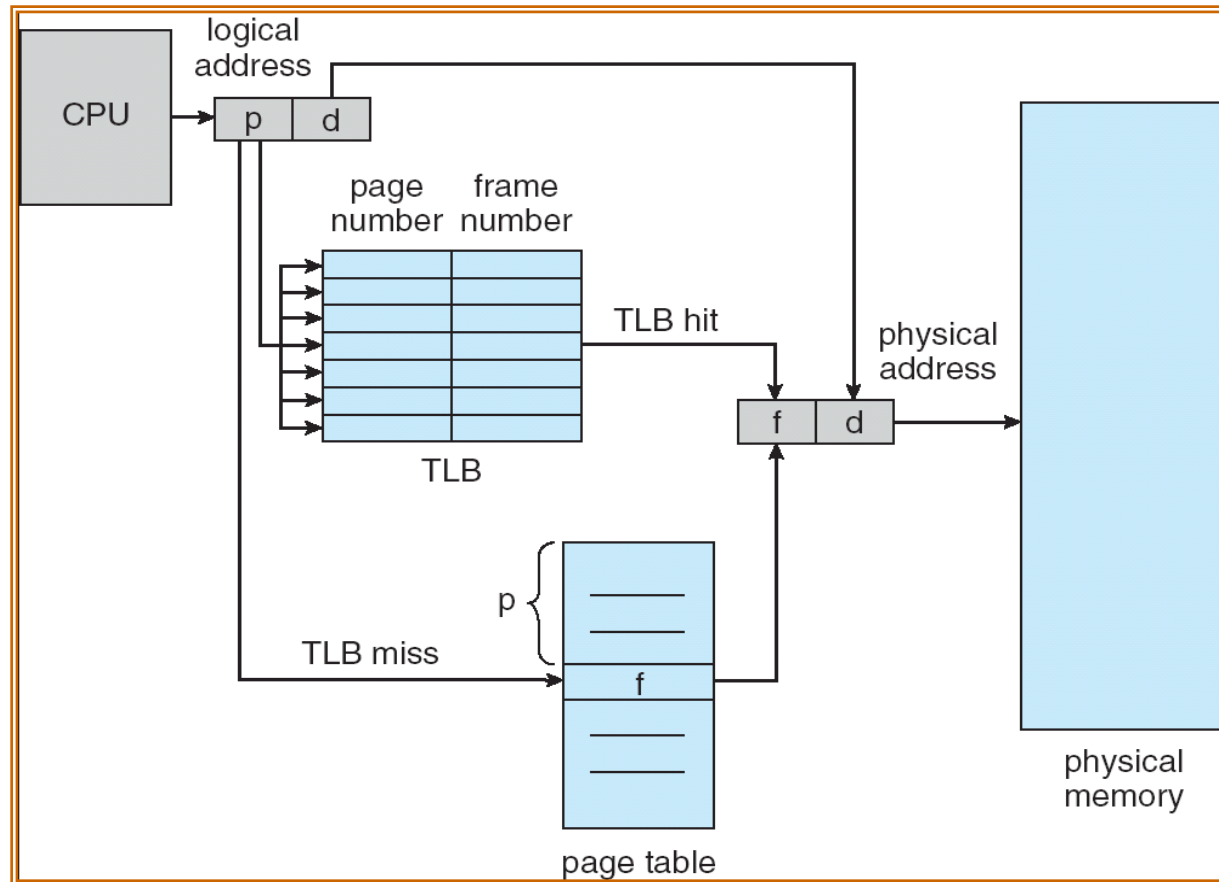
Virtual Memory

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PRLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses. One for the page table and one for the data/instruction.
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

Virtual Memory

Paging Hardware With TLB



Virtual Memory

TLBs – Translation Lookaside Buffers

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

A TLB to speed up paging

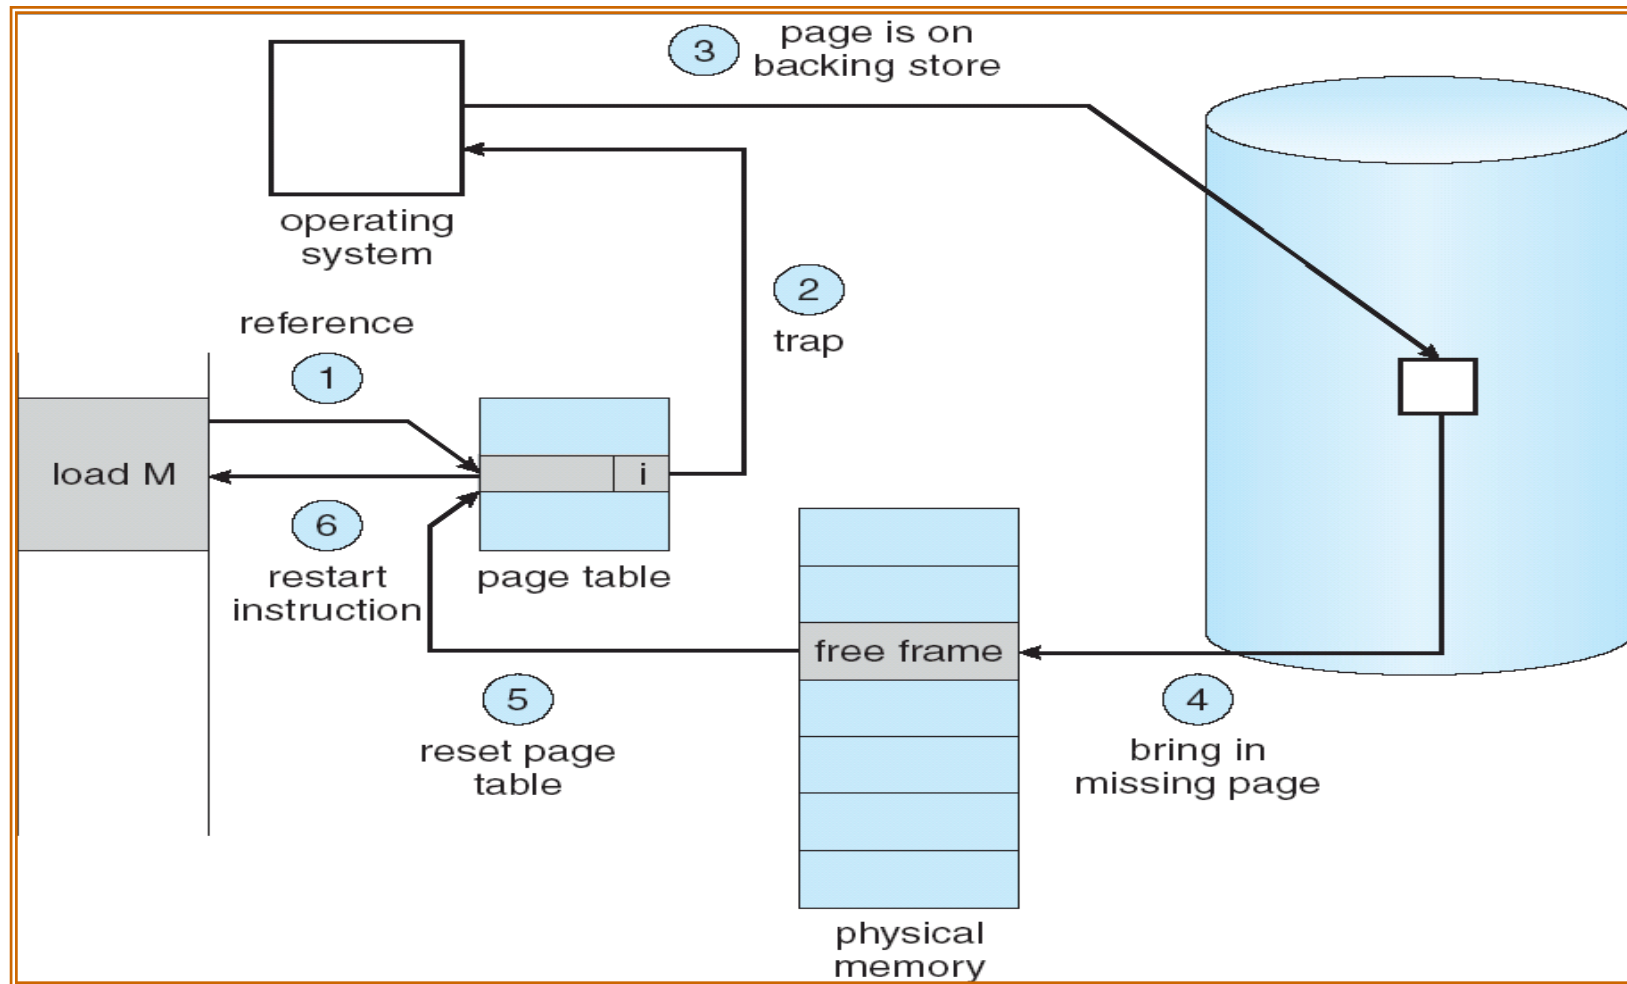
Virtual Memory

Page Fault

1. If there is a reference to a page, Just not in memory: **page fault**,
2. Trap to operating system:
3. Get empty page frame, Determine page on backing store
4. Swap page from disk into page frame in memory
5. Modifies page tables, Set validation bit = **v**
6. Restart the instruction that caused the page fault

Virtual Memory

Steps in Handling a Page Fault



Virtual Memory

Page Replacement

- What happens if there is no free frame?
- Page replacement – find some page in memory, but not really in use, swap it out
 - algorithm
 - performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times


Virtual Memory

Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

Virtual Memory

Page Replacement



3.4 Page Replacement Algorithms

Page Replacement Algorithms

- Optimal page replacement algorithm
- Not recently used page replacement
- First-In, First-Out page replacement
- Second chance page replacement
- Clock page replacement
- Least recently used page replacement
- Working set page replacement
- WSClock page replacement

Page Replacement Algorithms

Optimal Page Replacement Algorithm

- Replace page needed at the farthest point in future
 - Optimal but unrealizable
- Estimate by ...
 - logging page use on previous runs of process
 - although this is impractical

Page Replacement Algorithms

Not Recently Used Page Replacement Algorithm

- Each page has Reference bit, Modified bit
 - bits are set when page is referenced, modified
- Pages are classified
 1. not referenced, not modified
 2. not referenced, modified
 3. referenced, not modified
 4. referenced, modified
- NRU removes page at random
 - from lowest numbered non empty class

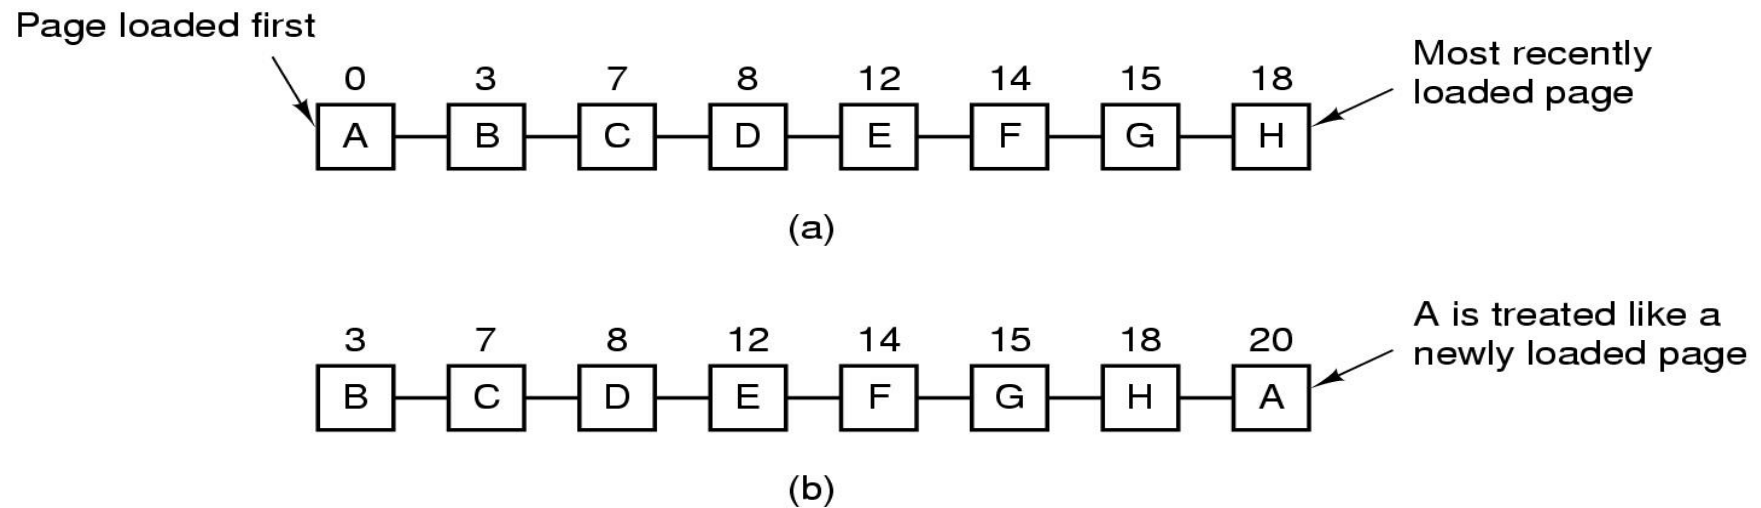
Page Replacement Algorithms

FIFO Page Replacement Algorithm

- Maintain a linked list of all pages
 - in order they came into memory
- Page at beginning of list replaced
- Disadvantage
 - page in memory the longest may be often used

Page Replacement Algorithms


Second Chance Page Replacement Algorithm



- Operation of a second chance
 - pages sorted in FIFO order
 - Page list if fault occurs at time 20, A has *R* bit set (numbers above pages are loading times)

Page Replacement Algorithms

The Clock Page Replacement Algorithm



When a page fault occurs, the page the hand is pointing to is inspected. The action taken depends on the R bit:

R = 0: Evict the page

R = 1: Clear R and advance hand

Page Replacement Algorithms

Least Recently Used (LRU)

- Assume pages used recently will be used again soon
 - throw out page that has been unused for longest time
- Must keep a linked list of pages
 - most recently used at front, least at rear
 - update this list every memory reference !!
- Alternatively keep counter in each page table entry
 - choose page with lowest value counter
 - periodically zero the counter

Page Replacement Algorithms

LRU

	Page			
	0	1	2	3
0	0	1	1	1
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0

(a)

	Page			
	0	1	2	3
0	0	0	1	1
1	1	0	1	1
2	0	0	0	0
3	0	0	0	0

(b)

	Page			
	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	1	1	0	1
3	0	0	0	0

(c)

	Page			
	0	1	2	3
0	0	0	0	0
1	1	0	0	0
2	1	1	0	0
3	1	1	1	0

(d)

	Page			
	0	1	2	3
0	0	0	0	0
1	1	0	0	0
2	1	1	0	1
3	1	1	0	0

(e)

	0	0	0	0
	1	0	1	1
	1	0	0	1
	1	0	0	0

(f)

	0	1	1	1
	0	0	1	1
	0	0	0	1
	0	0	0	0

(g)

	0	1	1	0
	0	0	1	0
	0	0	0	0
	1	1	1	0

(h)

	0	1	0	0
	0	0	0	0
	1	1	0	1
	1	1	0	0

(i)

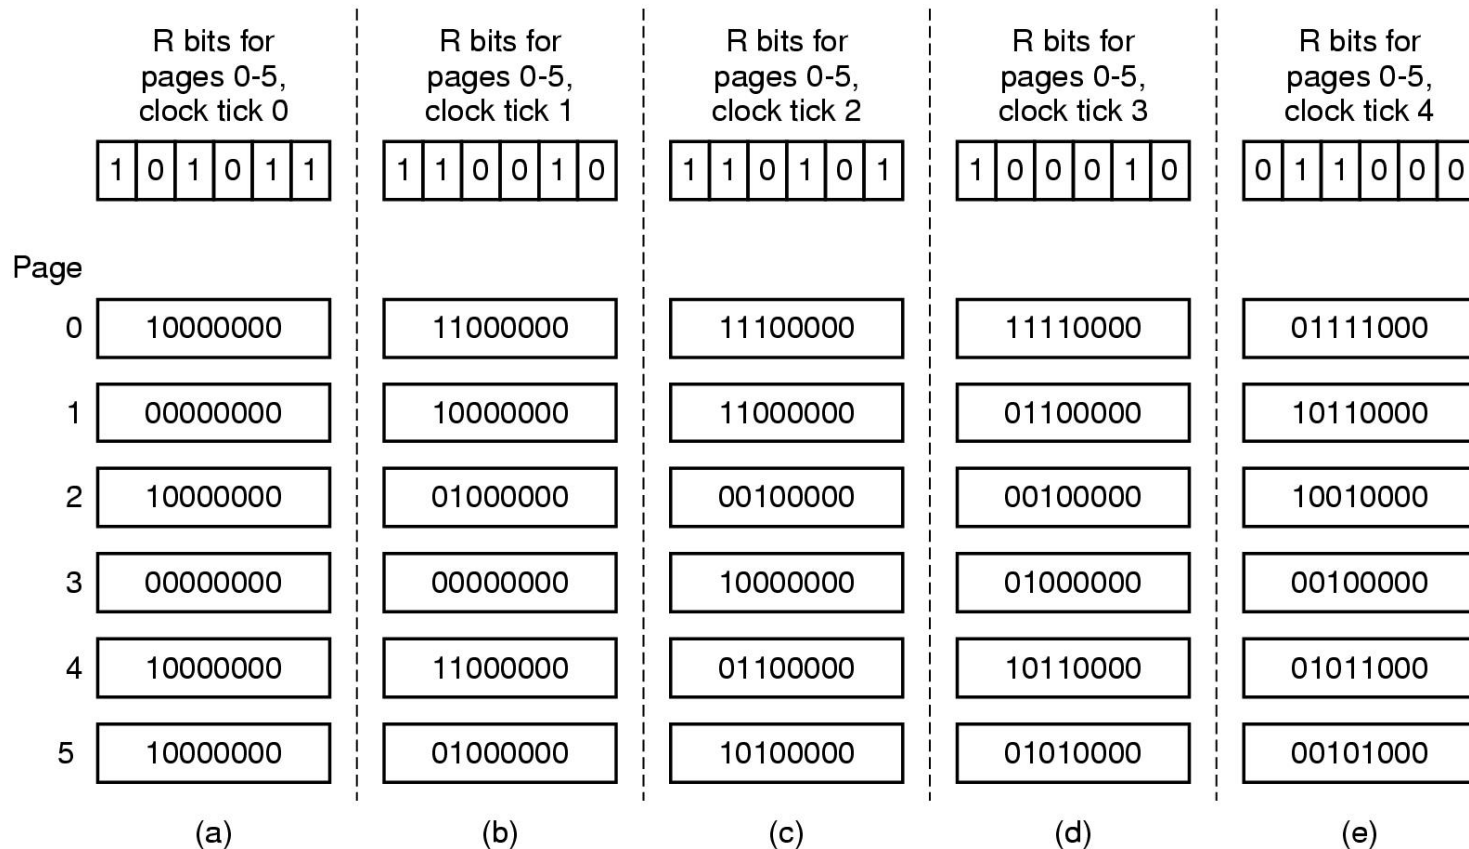
	0	1	0	0
	0	0	0	0
	1	1	0	0
	1	1	1	0

(j)

LRU using a matrix when pages are referenced in the order 0, 1, 2, 3, 2, 1, 0, 3, 2, 3.

Page Replacement Algorithms


Simulating LRU in Software



The aging algorithm simulates LRU in software. Shown are six page for five clock ticks. The five clock ticks are represented by (a) to (e).

Page Replacement Algorithms


Working Set Page Replacement (1)



- . The working set is the set of pages used by the k most recent memory references. The function $w(k, t)$ is the size of the working set at time t .

Page Replacement Algorithms

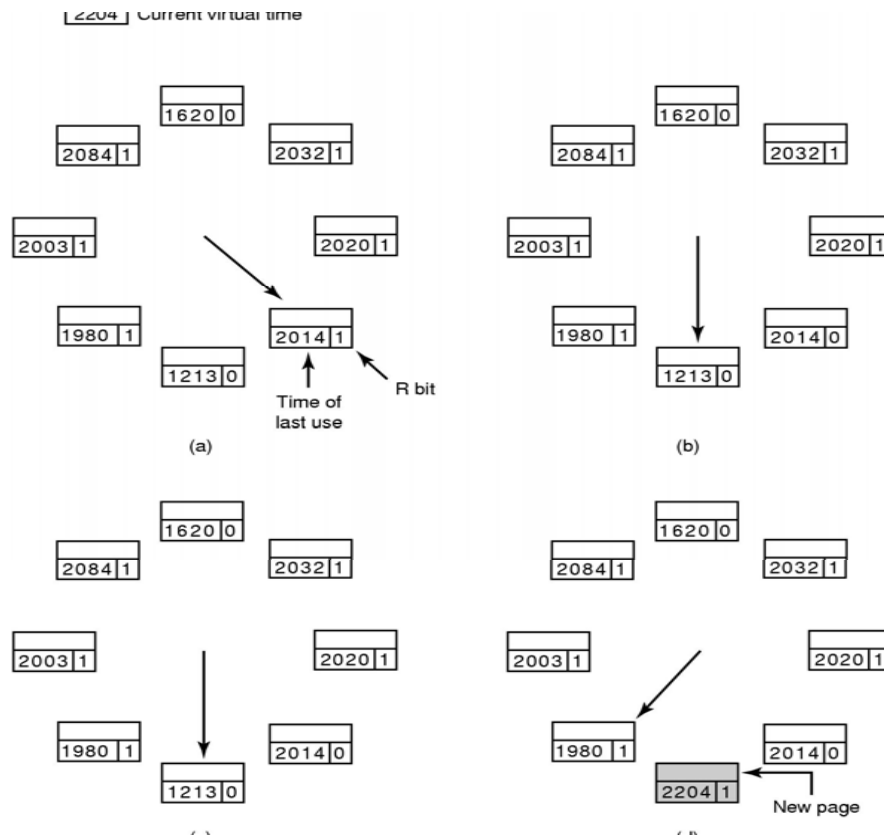
Working Set Page Replacement (2)



The working set algorithm.

Page Replacement Algorithms

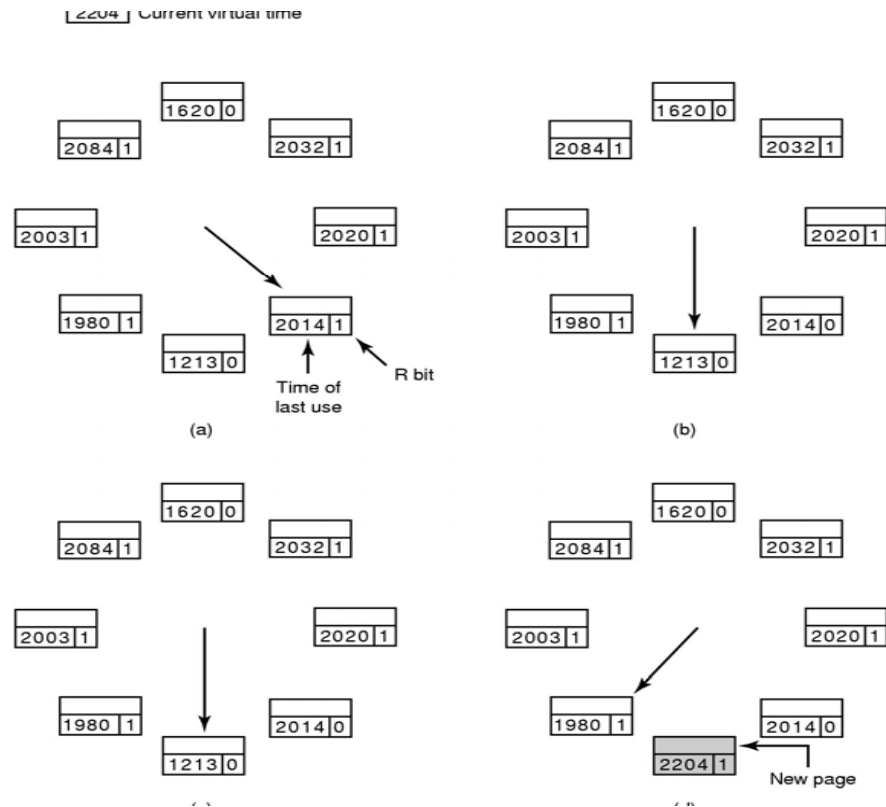
The WSClock Page Replacement Algorithm (3)



Operation of the WSClock algorithm. (a) and (b) give an example of what happens when $R = 1$.

Page Replacement Algorithms

The WSClock Page Replacement Algorithm (4)



Operation of the WSClock algorithm.
(c) and (d) give an example of $R = 0$.

Page Replacement Algorithms

The WSClock Page Replacement Algorithm (5)

When the hand comes all the way around to its starting point there are two cases to consider:

- At least one write has been scheduled.
- No writes have been scheduled.

Page Replacement Algorithms

Summary

Algorithm	Comment
Optimal	Not implementable, but useful as a benchmark
NRU (Not Recently Used)	Very crude
FIFO (First-In, First-Out)	Might throw out important pages
Second chance	Big improvement over FIFO
Clock	Realistic
LRU (Least Recently Used)	Excellent, but difficult to implement exactly
NFU (Not Frequently Used)	Fairly crude approximation to LRU
Aging	Efficient algorithm that approximates LRU well
Working set	Somewhat expensive to implement
WSClock	Good efficient algorithm

1.5 Design Issues For Paging System

Design Issues For Paging System

Local versus Global Allocation Policies (1)


	Age		
A0	10	A0	A0
A1	7	A1	A1
A2	5	A2	A2
A3	4	A3	A3
A4	6	A4	A4
A5	3	A6	A5
B0	9	B0	B0
B1	4	B1	B1
B2	6	B2	B2
B3	2	B3	A6
B4	5	B4	B4
B5	6	B5	B5
B6	12	B6	B6
C1	3	C1	C1
C2	5	C2	C2
C3	6	C3	C3

(a)
(b)
(c)

(a) Original configuration. (b) Local page replacement.
(c) Global page replacement.

Design Issues For Paging System


Local versus Global Allocation Policies (2)



Page fault rate as a function
of the number of page frames assigned.

Design Issues For Paging System


Separate Instruction and Data Spaces



- (a) One address space.
(b) Separate I and D spaces.

Design Issues For Paging System


Shared Pages



Two processes sharing the same program
sharing its page table.

Design Issues For Paging System

Shared Libraries



A shared library being used by two processes.

1.6 Implementation issues

Implementation Issues

Operating System Involvement with Paging

Four times when OS involved with paging

1. Process creation
 - determine program size
 - create page table
2. Process execution
 - MMU reset for new process
 - TLB flushed
3. Page fault time
 - determine virtual address causing fault
 - swap target page out, needed page in
4. Process termination time
 - release page table, pages

Implementation Issues

Page Fault Handling (1)

1. Hardware traps to kernel
2. General registers saved
3. OS determines which virtual page needed
4. OS checks validity of address, seeks page frame
5. If selected frame is dirty, write it to disk

Implementation Issues

Page Fault Handling (2)

6. OS brings schedules new page in from disk
7. Page tables updated
8. Faulting instruction backed up to when it began
9. Faulting process scheduled
10. Registers restored, Program continues


Implementation Issues

Locking Pages in Memory


- Virtual memory and I/O occasionally interact
- Proc issues call for read from device into buffer
 - while waiting for I/O, another processes starts up
 - has a page fault
 - buffer for the first proc may be chosen to be paged out
- Need to specify some pages locked
 - exempted from being target pages

Implementation Issues

Backing Store



(a)



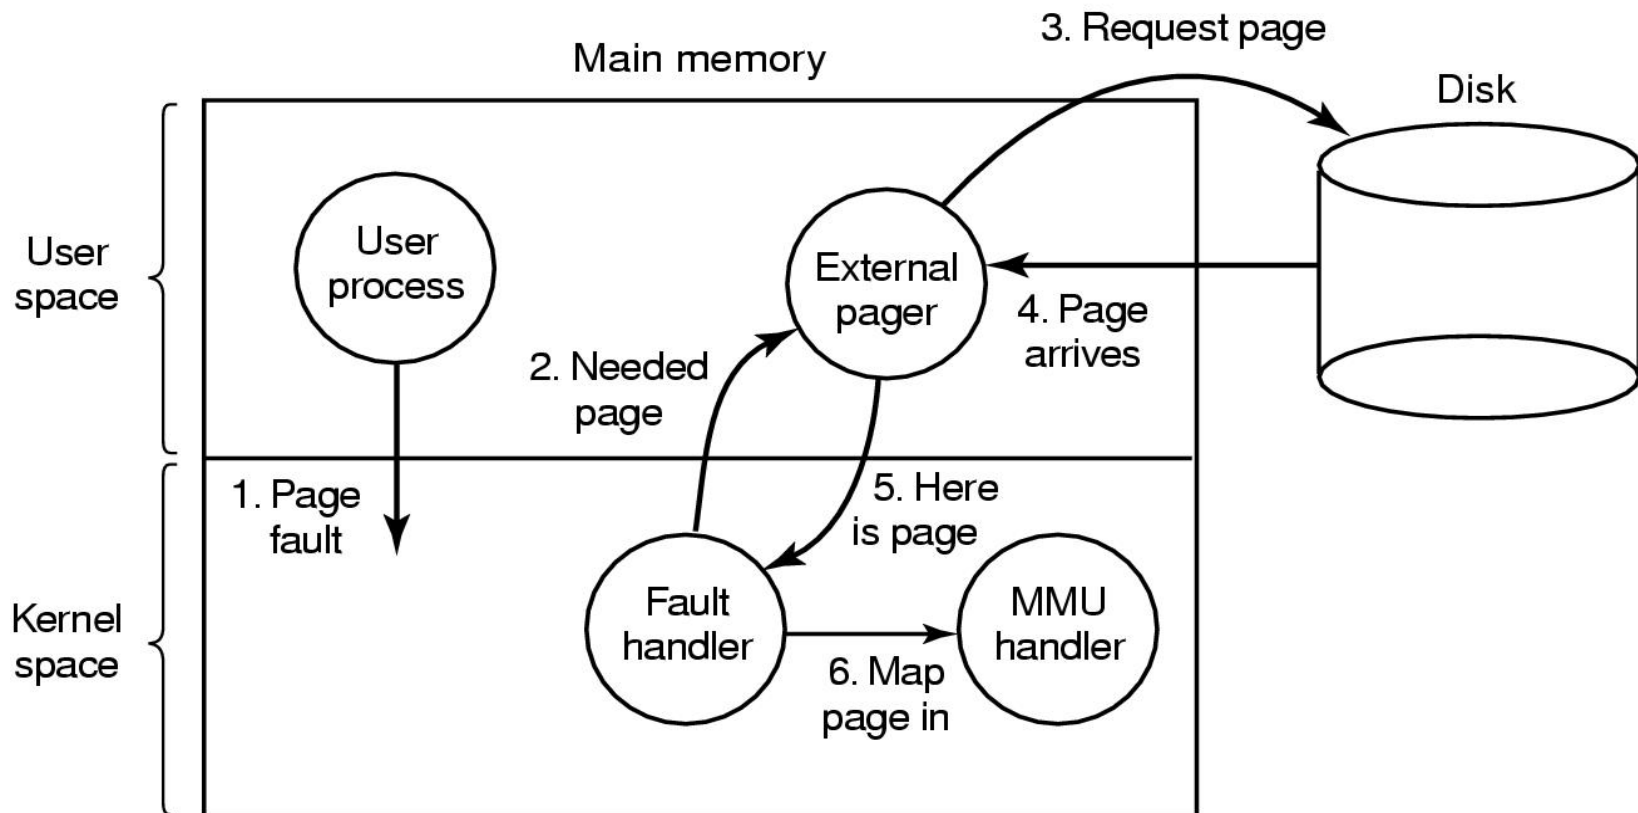
(b)

(a) Paging to static swap area

(b) Backing up pages dynamically

Implementation Issues

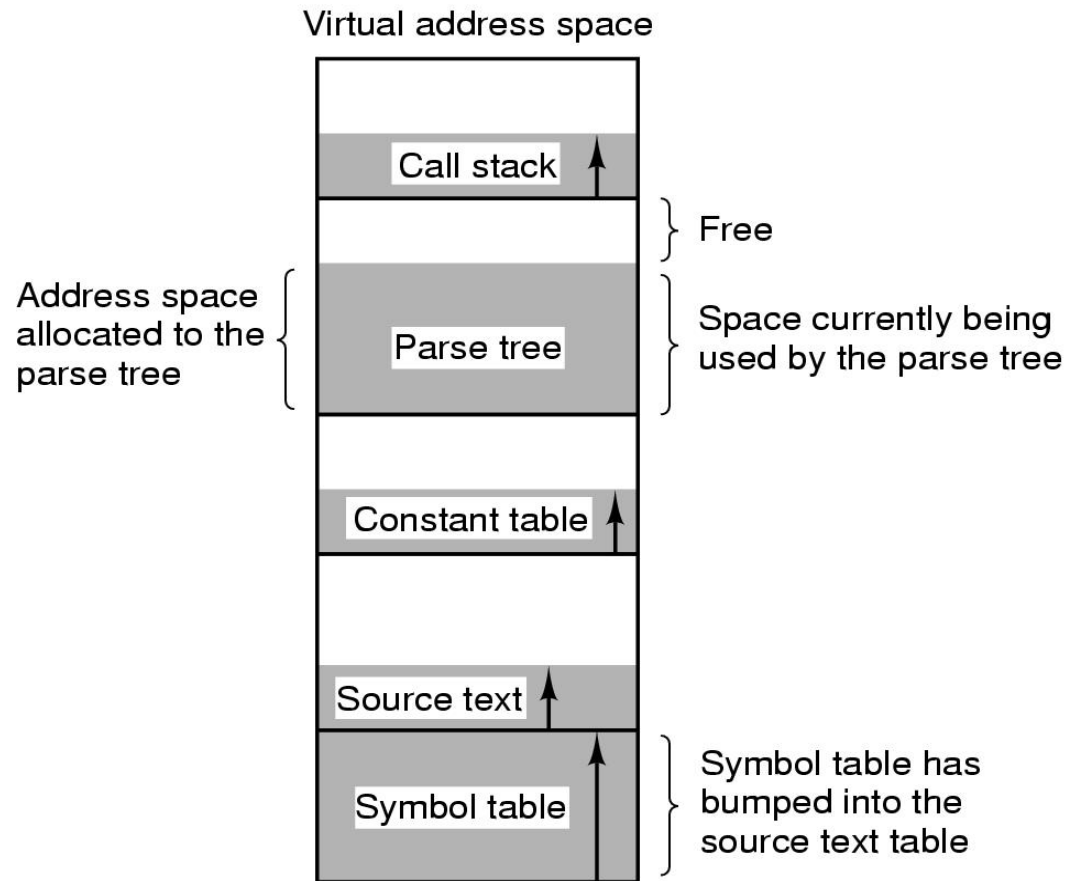
Separation of Policy and Mechanism



Page fault handling with an external pager

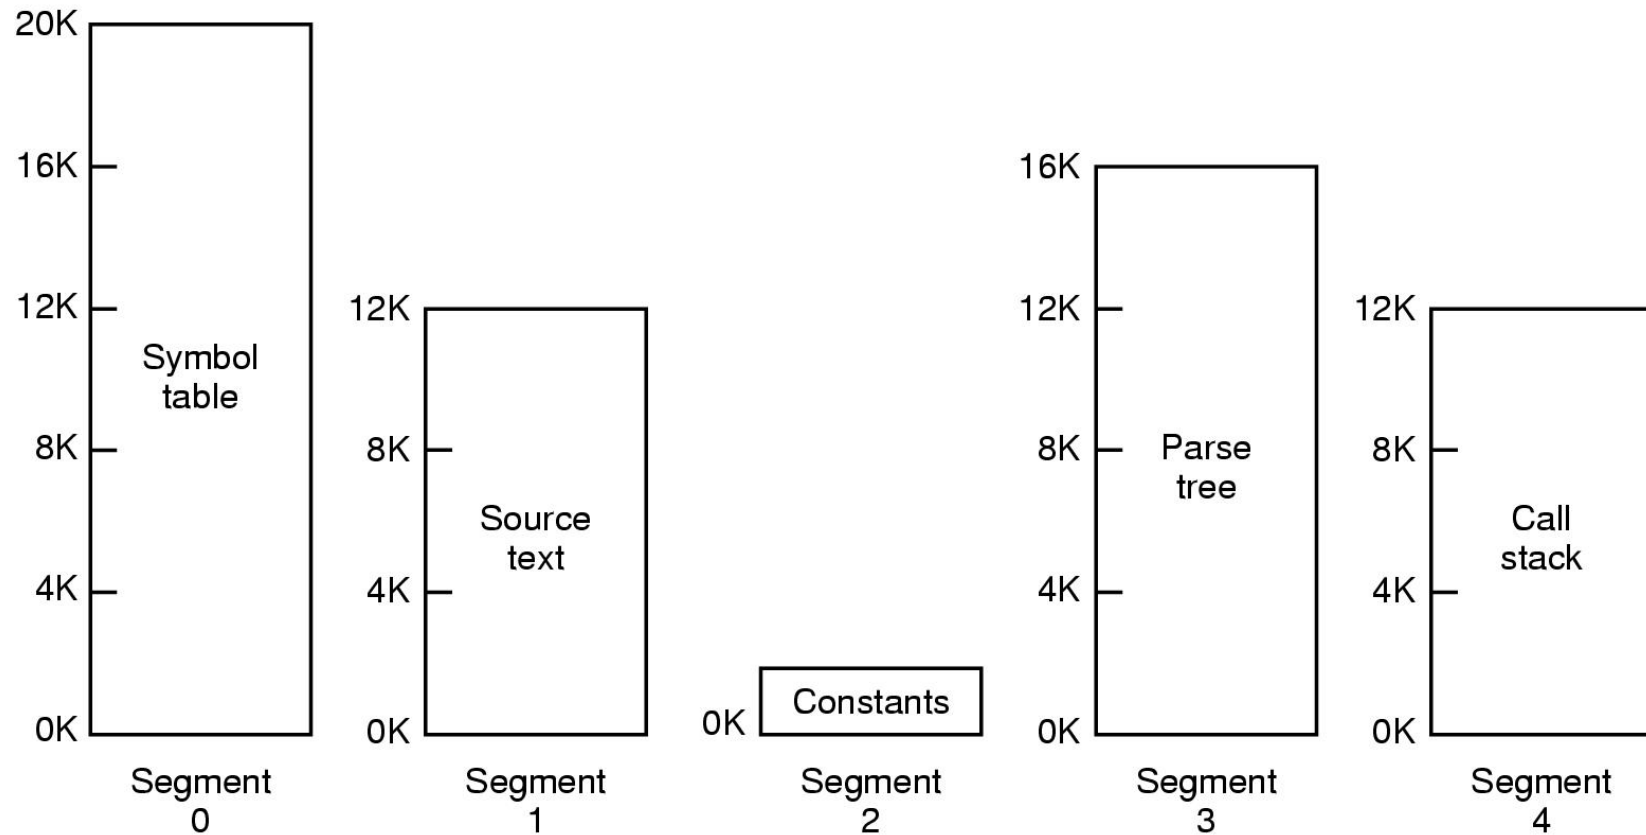
1.7 Virtual Memory Segmentation

Segmentation (1)



- One-dimensional address space with growing tables
- One table may bump into another

Segmentation (2)



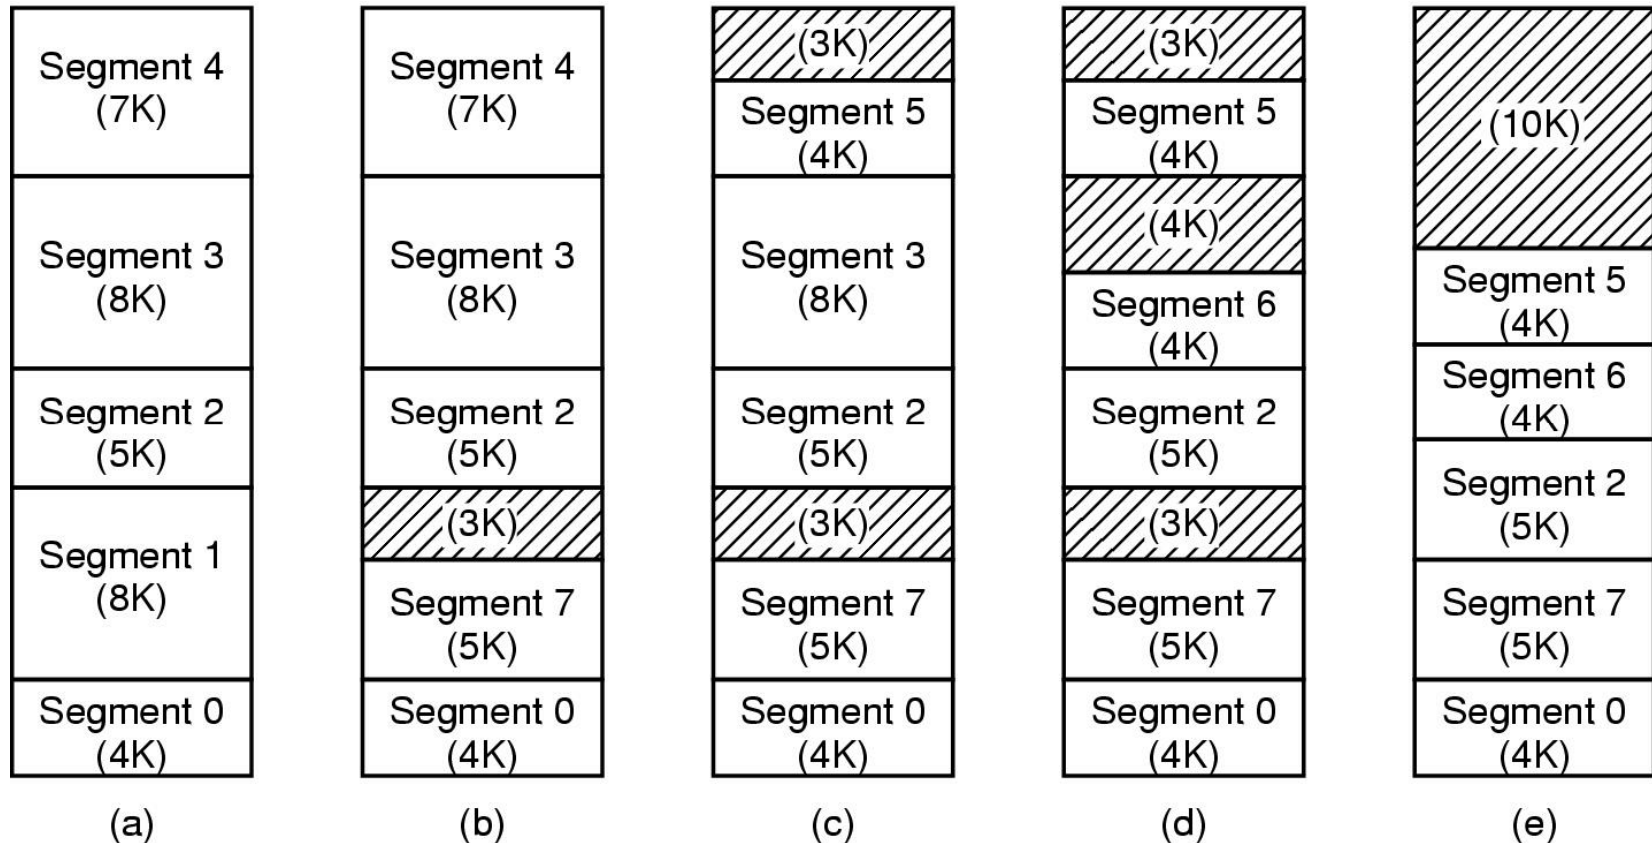
Allows each table to grow or shrink, independently

Segmentation (3)

Consideration	Paging	Segmentation
Need the programmer be aware that this technique is being used?	No	Yes
How many linear address spaces are there?	1	Many
Can the total address space exceed the size of physical memory?	Yes	Yes
Can procedures and data be distinguished and separately protected?	No	Yes
Can tables whose size fluctuates be accommodated easily?	No	Yes
Is sharing of procedures between users facilitated?	No	Yes
Why was this technique invented?	To get a large linear address space without having to buy more physical memory	To allow programs and data to be broken up into logically independent address spaces and to aid sharing and protection

Comparison of paging and segmentation

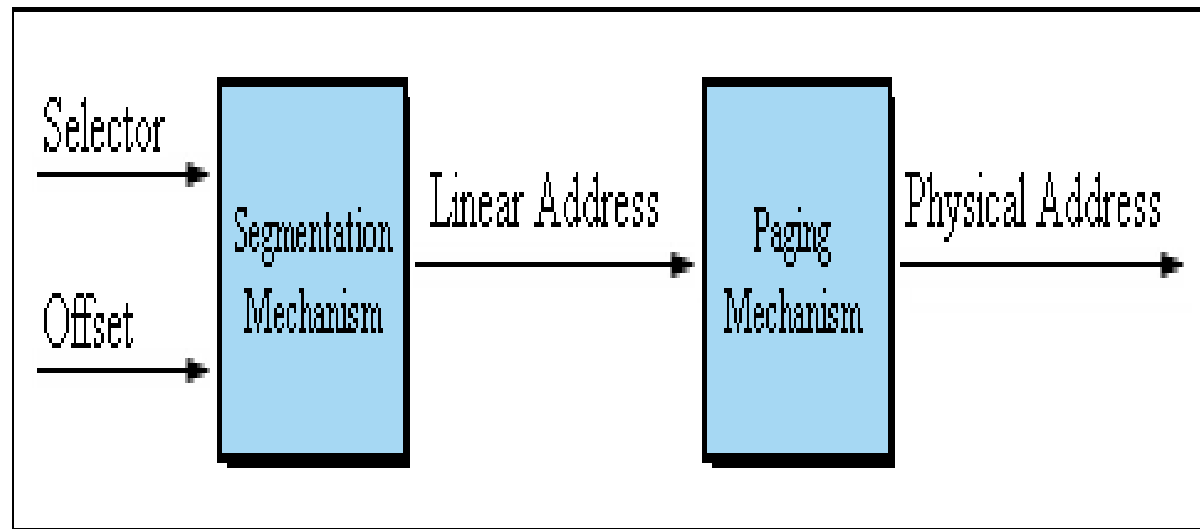
Implementation of Pure Segmentation (4)



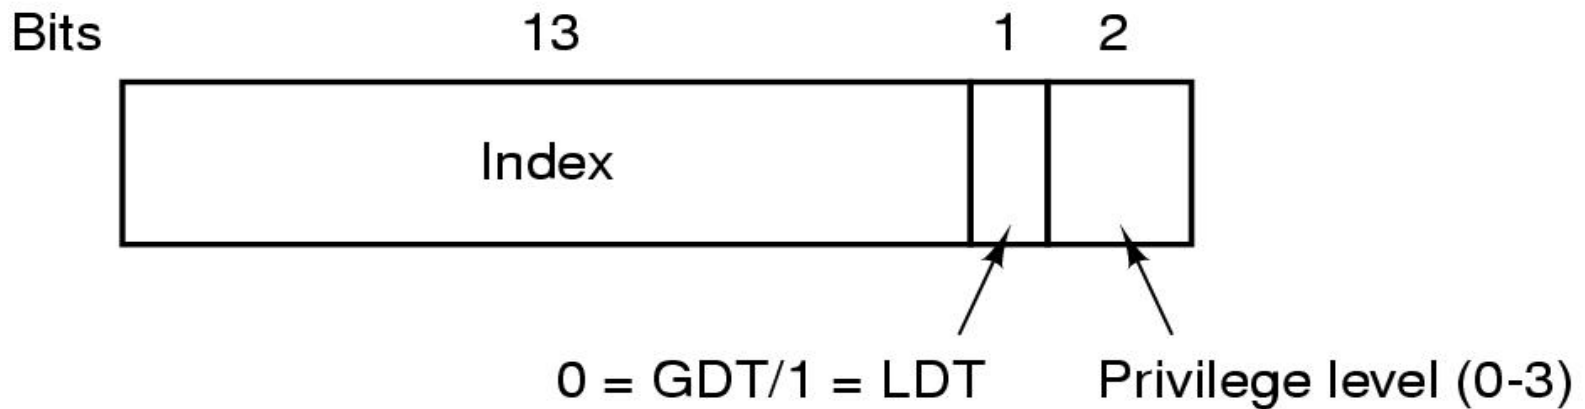
(a)-(d) Development of checkers

(e) Removal of the checkers by compaction

Segmentation with Paging: Pentium (1)



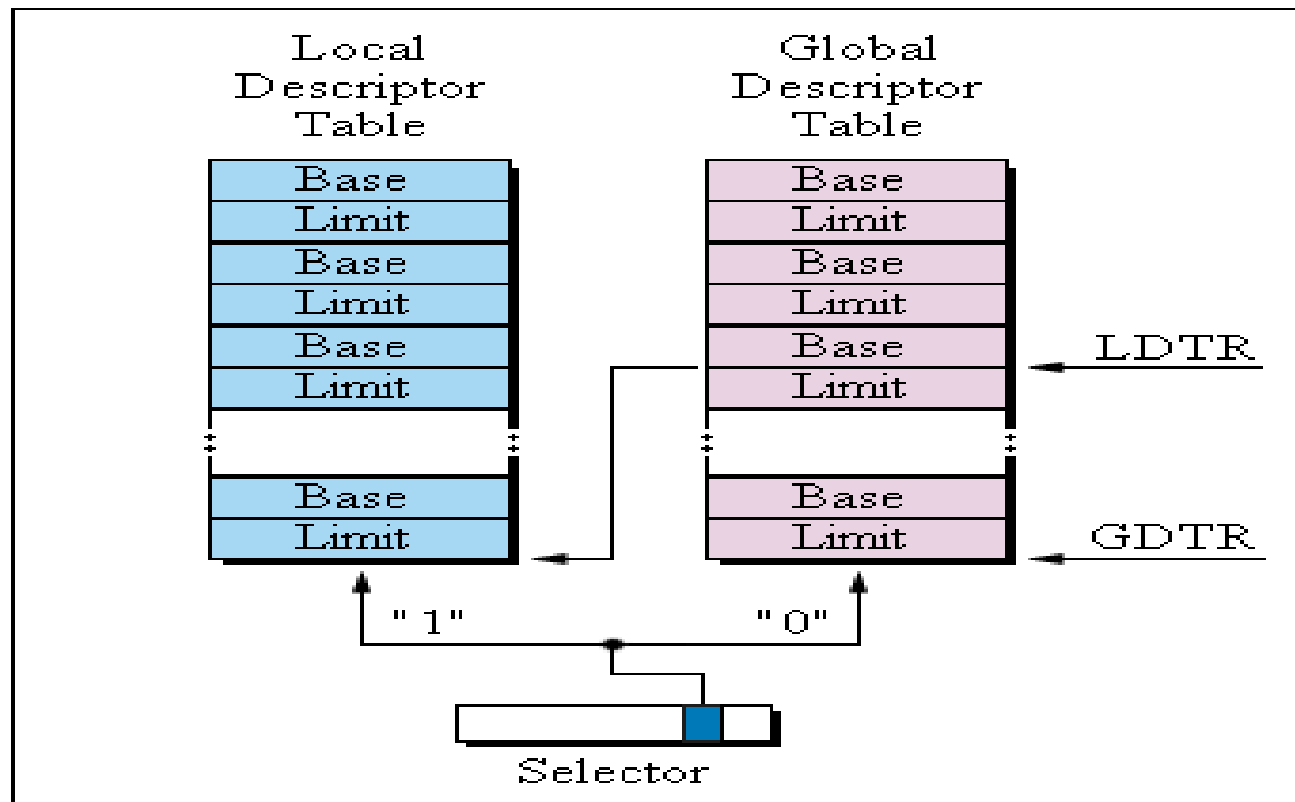
Segmentation with Paging: Pentium (2)



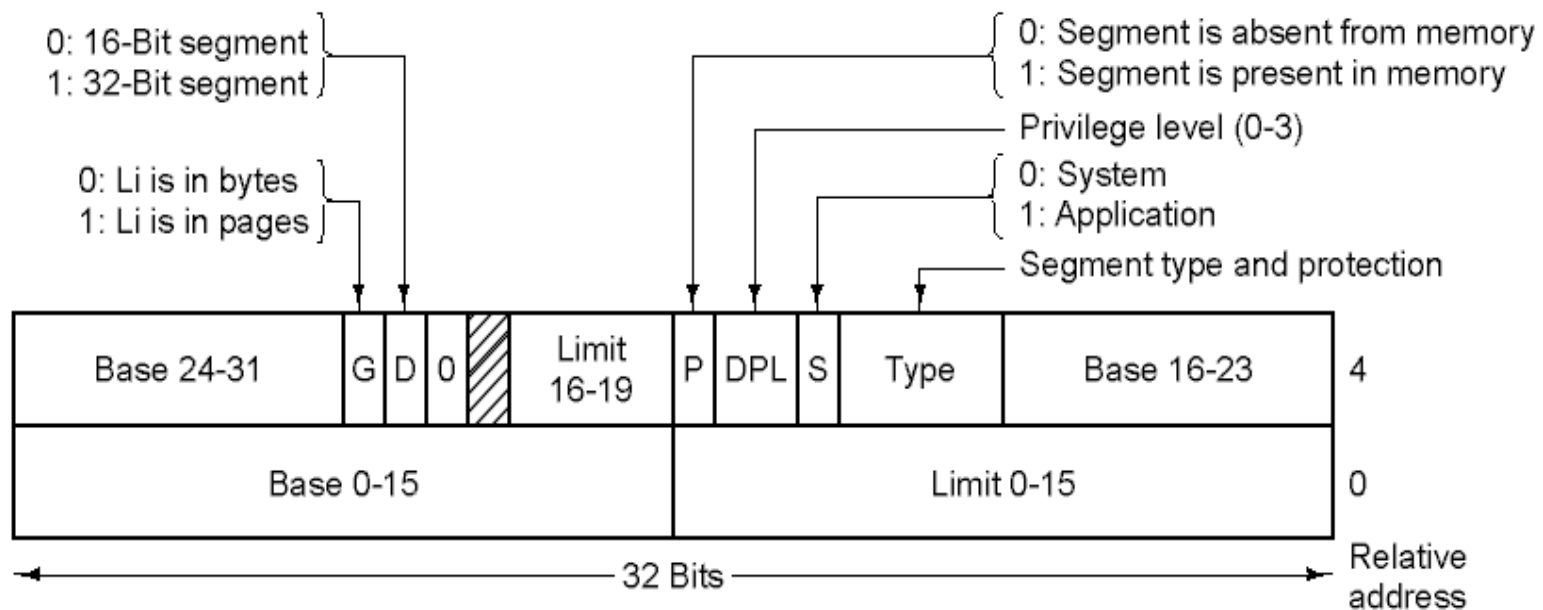
A Pentium selector

GDT (Global Descriptor Table), LDT (Local Descriptor Table)

Segmentation with Paging: Pentium (3)




Segmentation with Paging: Pentium (4)



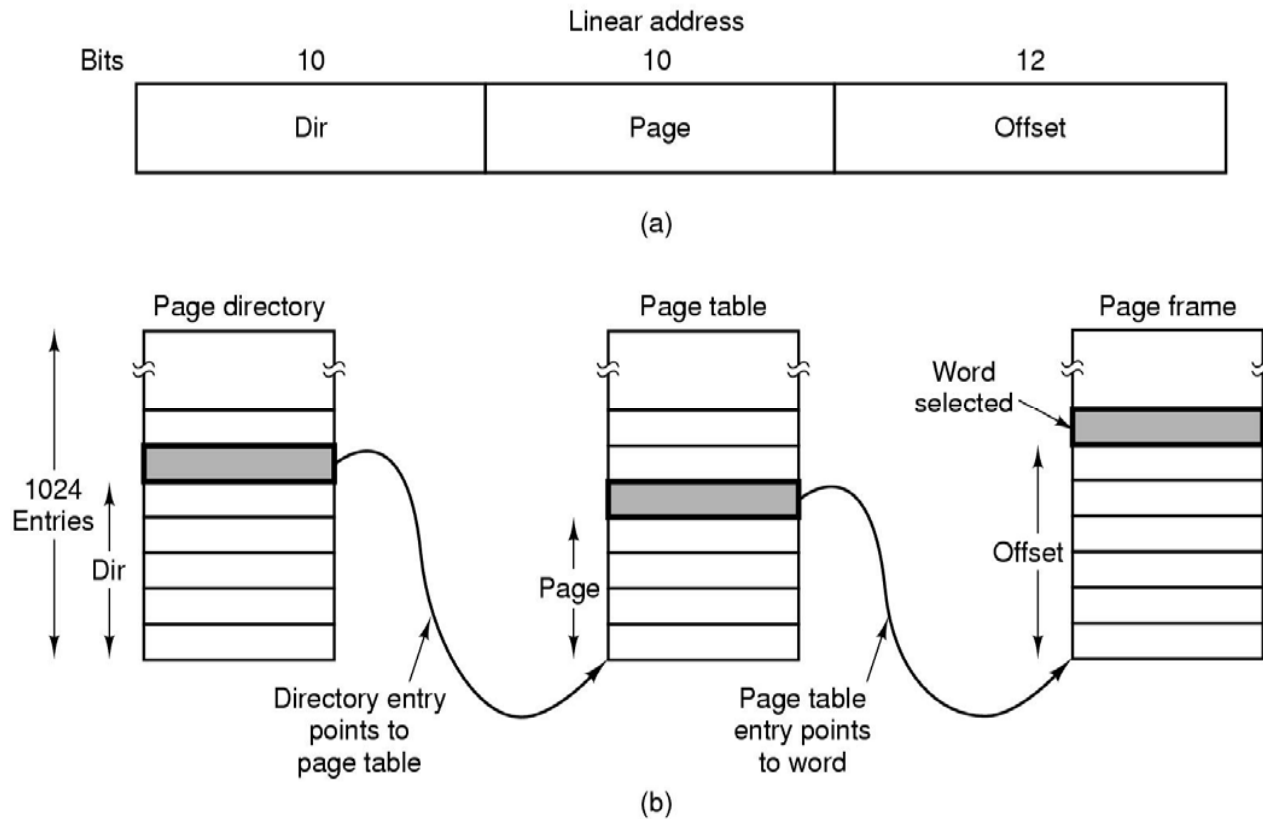
- Pentium code segment descriptor
- Data segments differ slightly

Segmentation with Paging: Pentium (5)



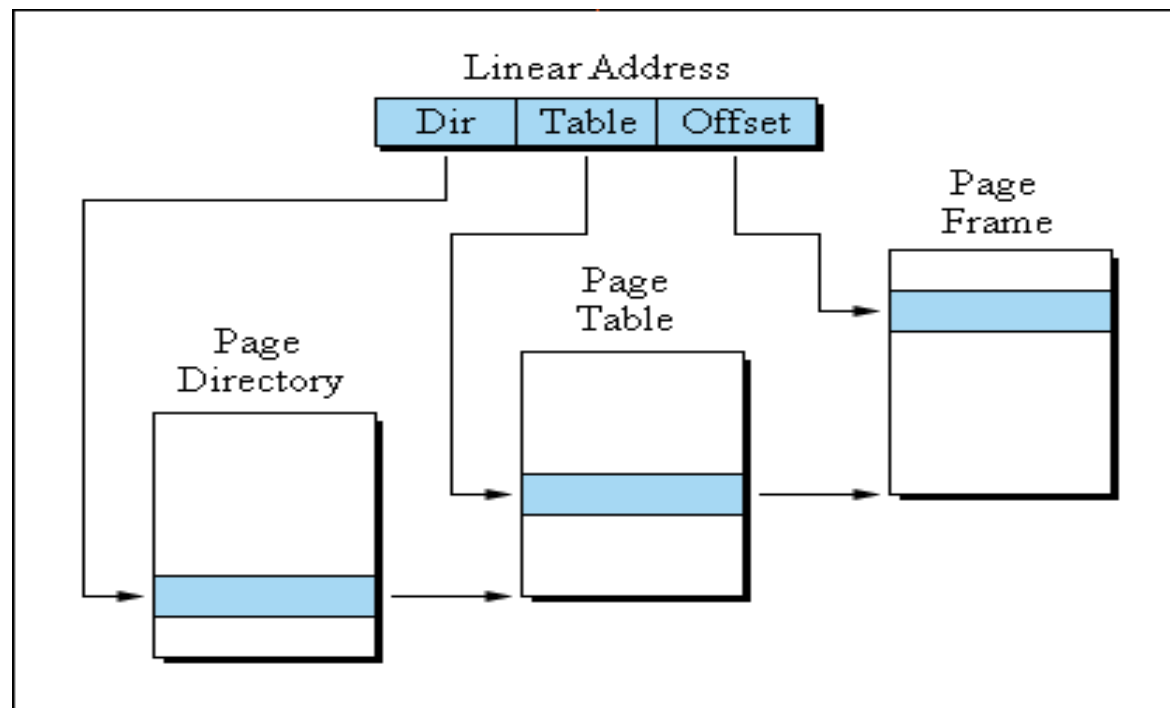
Conversion of a (selector, offset) pair to a linear address

Segmentation with Paging: Pentium (6)

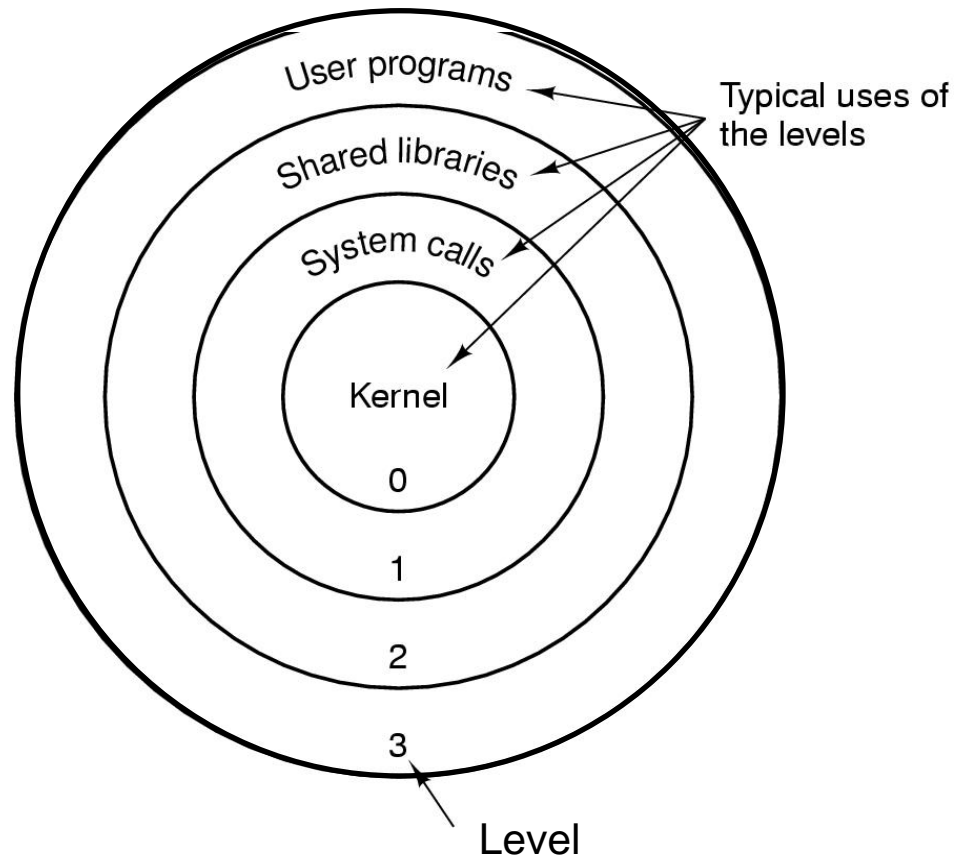


Mapping of a linear address onto a physical address

Segmentation with Paging: Pentium (7)



Segmentation with Paging: Pentium (8)



Protection on the Pentium